

This was the first python project for me, so it was extra hard. It took me more than 15 hours and some help from my friends to do it all.

I managed at the end to get 2.4 percent error on the validation set and 2.6 percent test error.

I have used 2 hidden layers, the first one has 128 neurons and a relu activation function. The second one has 96 neurons and a sigmoid activation function.

I tried to add a third layer but the results were poor.

The learning rate was 0.15 and max epochs is 30. The learning takes around 176.0s

